

Tweet Sentiment's Impact on Stock Returns的数据分析及预处理

学号: 1120202579 姓名: 彭高鹏

```
In [1]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

1. 数据集简介

该数据集包含 862,231 条带标签的推文和相关的股票回报，全面了解社交媒体对公司层面股票市场表现的影响。对于每条推文，研究人员都提取了推文日期及其相关股票代码等数据，以及最新价格和各种回报（1 天回报、2 天回报、3 天回报、7 天回报）等指标。还记录了 10 天间隔和 30 天间隔的波动率分数。最后，包括来自长短期记忆（LSTM）和 TextBlob 模型的情绪分数，以量化这些消息传递的整体基调。

TWEET:推文的文本 (String)

STOCK:推文中提到的公司股票 (String)

DATE:推文的发布日期 (Date)

LAST_PRICE:公司发布推文时的最后价格 (Float)

1_DAY_RETURN:股票在推特上发布后的第二天退回或损失的金额 (Float)

2_DAY_RETURN:在推文发布后的两天内退回或损失的股票金额 (Float)

3_DAY_RETURN:股票在被推文发布后的三天内退回或损失的金额 (Float)

7_DAY_RETURN:在推文发布后的 7 天内退回或损失的股票金额 (Float)

PX_VOLUME:发推文时的交易量 (Integer)

VOLATILITY_10D:跨 10 天窗口的波动率衡量 (Float)

VOLATILITY_30D:波动率衡量 30 天窗口 (Float)

LSTM_POLARITY来自 LSTM 的标记情绪 (Float)

TEXTBLOB_POLARITY:来自 TextBlob 的标记情绪 (Float)

下面对数据集的csv文件进行读取，并展示前几个数据对象

```
In [2]: df = pd.read_csv('full_dataset-release.csv')
df.info()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_25684\3661070278.py:1: DtypeWarning: Columns (0) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv('full_dataset-release.csv')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1102013 non-null object
1   TWEET                 1395398 non-null object
2   STOCK                862231 non-null object
3   DATE                 862231 non-null object
4   LAST_PRICE           862231 non-null float64
5   1_DAY_RETURN         862231 non-null float64
6   2_DAY_RETURN         862231 non-null float64
7   3_DAY_RETURN         862231 non-null float64
8   7_DAY_RETURN         862231 non-null float64
9   PX_VOLUME            862231 non-null float64
10  VOLATILITY_10D       862130 non-null float64
11  VOLATILITY_30D       862184 non-null float64
12  LSTM_POLARITY        661102 non-null float64
13  TEXTBLOB_POLARITY    367665 non-null float64
dtypes: float64(10), object(4)
memory usage: 149.1+ MB
```

```
In [3]: df.head(5)
```

Out[3]:

	Unnamed: 0	TWEET	STOCK	DATE	LAST_PRICE	1_DAY_RETURN	2_DA
0	0	RT @robertoglezcano: @amazon #Patents Show Fl...	NaN	NaN	NaN	NaN	
1	NaN	Amazon	31/01/2017	823.48	0.008379	0.014924	
2	1	@FAME95FM1 Jamaicans make money with @Payoneer...	PayPal	31/01/2017	39.780000	0.002011	
3	2	@CBSi Jamaicans make money with @Payoneer @Pay...	PayPal	31/01/2017	39.780000	0.002011	
4	3	@Hitz92fm Jamaicans make money with @Payoneer ...	PayPal	31/01/2017	39.780000	0.002011	

定义分析数据的方案，包括展示数据信息、绘制数据的直方图和盒图

```
In [4]: #定义标称型数据的分析方案
def analyze_freq_nominal(name, data, draw=True):
    freq = data.value_counts()
    print("Frequency information of the {}: \n".format(freq))
    print("Missing Value Count:", data.isnull().sum())

    if draw:
        # 绘制处理数据的频数直方图
        plt.figure(figsize=(12, 5))
        plt.bar(freq.index, freq.values)
        plt.title(name + ' Frequency Histogram')
        plt.xlabel(name)
        plt.ylabel(' Frequency')
        plt.grid(axis='y')
        plt.tick_params(axis='x', labelsiz=8)
        plt.xticks(rotation=-50)
        plt.tight_layout()
        plt.show()

#定义数值型数据的分析方案
def analyze_freq_and_box_numeric(name, data, draw=True, bin_set=None):
    print("The Information of {}: \n".format(name))
    print(data.describe())
    print("Missing Value Count:", data.isnull().sum())

    if draw:
        # 绘制数据的频数直方图
        plt.figure(figsize=(24, 5))
        plt.subplot(1, 2, 1)
        if bin_set:
            sns.histplot(data, bins=bin_set, kde=True)
        else:
            sns.histplot(data, kde=True)
        plt.title(name + ' Frequency Histogram')
        plt.xlabel(name)
        plt.ylabel(' Frequency')
        plt.grid(axis='y')

        #绘制数据的盒图
        plt.subplot(1, 2, 2)
        sns.boxplot(data)
        plt.title(name + ' Boxplot')
        plt.ylabel(name)
        plt.grid(axis='y')

        plt.tight_layout()
        plt.show()
```

2. 数据摘要和可视化

• 数据摘要

标称属性，给出每个可能取值的频数

数值属性，给出5数概括及缺失值的个数

• 数据可视化

使用直方图、盒图等检查数据分布及离群点

2.1 分析标称型数据

数据集中的标称型数据有：推文的文本TWEET,推文中提到的公司股票STOCK
其中对于数据类型较多的数据，我们不进行画图，只进行频数分析

```
In [5]: nominal_df = df.select_dtypes(include=['object'])  
print(nominal_df.columns)
```

```
Index(['Unnamed: 0', 'TWEET', 'STOCK', 'DATE'], dtype='object')
```



数据集中的数值型数据有: LAST_PRICE, 1_DAY_RETURN, 2_DAY_RETURN, 3_DAY_RETURN, PX_VOLUME, VOLATILITY_10D, VOLATILITY_30D,

```
In [7]: numeric_df = df.select_dtypes(include=['int64', 'float64'])  
        print(numeric_df.columns)
```

```
Index(['LAST_PRICE', '1_DAY_RETURN', '2_DAY_RETURN', '3_DAY_RETURN',  
       '7_DAY_RETURN', 'PX_VOLUME', 'VOLATILITY_10D', 'VOLATILITY_30D',  
       'LSTM_POLARITY', 'TEXTBLOB_POLARITY'],  
      dtype='object')
```


The Information of LAST_PRICE:

```
count      862231.000000
mean         716.123477
std        4731.498118
min         -0.173554
25%          0.000000
50%          0.014220
75%         115.600000
max       165500.000000
Name: LAST_PRICE, dtype: float64
Missing Value Count: 533219
```

[illegible]

The Information of 1_DAY_RETURN:

```
count      862231.000000
mean         0.001473
std         0.023068
min        -0.177851
25%        -0.007511
50%         0.000000
75%         0.008192
max         0.243639
Name: 1_DAY_RETURN, dtype: float64
Missing Value Count: 533219
```



The Information of 2_DAY_RETURN:

The Information of 3 DAY RETURN:

Figure 1 consists of two plots side-by-side. The left plot is a frequency histogram titled '3_DAY_RETURN Frequency Histogram'. The x-axis is labeled '3 DAY RETURN' and ranges from 0.0 to 3.0 with a multiplier of $1e8$. The y-axis is labeled 'Frequency' and ranges from 0.0 to 2.5. The histogram shows a very high frequency (around 2.4) for returns near 0, with a long tail extending to the right. The right plot is a boxplot titled '3_DAY_RETURN Boxplot'. The x-axis is labeled '3 DAY RETURN' and ranges from 0.0 to 3.0 with a multiplier of $1e8$. The y-axis is labeled '3 DAY RETURN' and ranges from 0.0 to 3.0 with a multiplier of $1e8$. The boxplot shows a median near 0, with a long right tail and several outliers.

The Information of 7 DAY RETURN:

```
count      8.622310e+05
mean       6.201036e+06
std        1.388221e+07
min       -2.049587e-01
25%        5.474453e-03
50%        1.575000e+01
75%        4.766038e+06
max        1.562074e+08
Name: 7_DAY_RETURN, dtype: float64
Missing Value Count: 533219
```




```
count      8.622310e+05
mean       7.522187e+06
std        1.591847e+07
min        1.000000e+00
25%        2.051700e+01
50%        3.979400e+01
75%        6.531761e+06
max         3.148332e+08
Name: PX_VOLUME, dtype: float64
Missing Value Count: 533219
```



```
count      862130.000000
mean        18.293836
std         16.277630
min         -1.000000
25%         8.792000
50%        17.897000
75%        23.916000
max        124.137000
Name: VOLATILITY_10D, dtype: float64
Missing Value Count: 533320
```



The Information of VOLATILITY_30D:


Figure 1 consists of two plots. The left plot is a frequency histogram titled 'VOLATILITY_30D Frequency Histogram'. The x-axis is labeled 'VOLATILITY_30D' and ranges from 0 to 80. The y-axis is labeled 'frequency' and ranges from 0 to 300,000. The histogram shows a very high frequency of values near 0, with a long right tail extending to 80. The right plot is a boxplot titled 'VOLATILITY_30D Boxplot'. The y-axis is labeled 'VOLATILITY_30D' and ranges from 0 to 80. The boxplot shows a median around 20, with a long right tail and several outliers above 60.

The Information of LSTM_POLARITY:

The left plot, titled "LSTM_POLARITY Frequency Histogram", shows the frequency distribution of LSTM_POLARITY values. The x-axis ranges from -1.00 to 1.00, and the y-axis (Frequency) ranges from 0 to 200,000. The histogram bars are blue, and a blue normal distribution curve is overlaid, peaking around 0.00. The right plot, titled "LSTM_POLARITY Boxplot", shows the distribution of LSTM_POLARITY values. The y-axis ranges from -1.00 to 1.00. The box is blue, with a median line at 0.00. The whiskers extend from approximately -0.9 to 0.9.

The Information of TEXTBLOB POLARITY:

10/32



3. 数据缺失的处理

观察数据集中缺失数据，分析其缺失的原因。分别使用下列四种策略对缺失值进行处理：

- 将缺失部分剔除
- 用最高频率值来填补缺失值
- 通过属性的相关关系来填补缺失值
- 通过数据对象之间的相似性来填补缺失值

3.1 分析数据缺失的原因

该数据集几乎所有的属性都有或多或少数量的缺失值，从前面对标称型数据和数值型数据的分析来看，二者也各自都有部分缺失值。数据缺失的原因可能有以下几点：

- 股价信息的缺失，股票本身的信息不全
- 爬取股价信息的程序算法不完备

3.2 剔除有缺失值的数据对象

```
In [9]: data = df.dropna()
```

展示填补前的数据信息

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1102013 non-null object
1   TWEET                 1395398 non-null object
2   STOCK                862231 non-null object
3   DATE                 862231 non-null object
4   LAST_PRICE           862231 non-null float64
5   1_DAY_RETURN         862231 non-null float64
6   2_DAY_RETURN         862231 non-null float64
7   3_DAY_RETURN         862231 non-null float64
8   7_DAY_RETURN         862231 non-null float64
9   PX_VOLUME            862231 non-null float64
10  VOLATILITY_10D        862130 non-null float64
11  VOLATILITY_30D        862184 non-null float64
12  LSTM_POLARITY         661102 non-null float64
13  TEXTBLOB_POLARITY     367665 non-null float64
dtypes: float64(10), object(4)
memory usage: 149.1+ MB
```

In [11]:

df.head(10)

Out[11]:

	Unnamed: 0	TWEET	STOCK	DATE	LAST_PRICE	1_DAY_RETURN	2_DA
0	0	RT @robertoglezcano: @amazon #Patents Show Fl...	NaN	NaN	NaN	NaN	
1	NaN	Amazon	31/01/2017	823.48	0.008379	0.014924	
2	1	@FAME95FM1 Jamaicans make money with @Payoneer...	PayPal	31/01/2017	39.780000	0.002011	
3	2	@CBSi Jamaicans make money with @Payoneer @Pay...	PayPal	31/01/2017	39.780000	0.002011	
4	3	@Hitz92fm Jamaicans make money with @Payoneer ...	PayPal	31/01/2017	39.780000	0.002011	
5	4	RT @loadsofvans: Retweet this post & follo...	NaN	NaN	NaN	NaN	
6	NaN	Amazon	31/01/2017	823.48	0.008379	0.014924	
7	5	RT @Disney: See The Newest Trailer for #Beauty...	NaN	NaN	NaN	NaN	
8	NaN	Disney	31/01/2017	110.65	0.002621	-0.012201	
9	6	RT @tarastrong: Just got @facebook back! I REA...	NaN	NaN	NaN	NaN	

填补后的数据集信息如下：

```
In [12]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 367620 entries, 2 to 1395443
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             367620 non-null object
1   TWEET                  367620 non-null object
2   STOCK                  367620 non-null object
3   DATE                   367620 non-null object
4   LAST_PRICE             367620 non-null float64
5   1_DAY_RETURN           367620 non-null float64
6   2_DAY_RETURN           367620 non-null float64
7   3_DAY_RETURN           367620 non-null float64
8   7_DAY_RETURN           367620 non-null float64
9   PX_VOLUME              367620 non-null float64
10  VOLATILITY_10D          367620 non-null float64
11  VOLATILITY_30D          367620 non-null float64
12  LSTM_POLARITY           367620 non-null float64
13  TEXTBLOB_POLARITY       367620 non-null float64
dtypes: float64(10), object(4)
memory usage: 42.1+ MB
```

In [13]:

data.head(10)

Out[13]:

	Unnamed: 0	TWEET	STOCK	DATE	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN
2	1	@FAME95FM1 Jamaicans make money with @Payoneer...	PayPal	31/01/2017	39.78	0.002011	
3	2	@CBSi Jamaicans make money with @Payoneer @Pay...	PayPal	31/01/2017	39.78	0.002011	
4	3	@Hitz92fm Jamaicans make money with @Payoneer ...	PayPal	31/01/2017	39.78	0.002011	
11	7	RT @nikitakhara: Thank you, @Starbucks CEO for...	Starbucks	31/01/2017	55.22	0.012314	
20	12	@gawker Jamaicans make money with @Payoneer @P...	PayPal	31/01/2017	39.78	0.002011	
23	14	RT @cultcommoncore: Dumping @Starbucks\r\r\r\r\r...	Starbucks	31/01/2017	55.22	0.012314	
26	16	@amazon has your back when it comes to food sh...	Amazon	31/01/2017	823.48	0.008379	
35	21	RT @nia4_trump: So instead of hiring 10,000 un...	Starbucks	31/01/2017	55.22	0.012314	
48	28	Hmmm interesting... \r\r\r\r\r\r\r\r\r\r\nonly @IBM ha...	Disney	31/01/2017	110.65	0.002621	
53	31	RT @IndiaHistorypic: 1994::Young @sundarpichai...	Google	31/01/2017	820.19	0.004438	

比较处理前后数据集差异

下面以1_DAY_RETURN 这一属性为例，对比数据集在剔除含缺失值数据对象后的差异


```
In [15]: data = df.copy(deep=True)
```

展示填补前的数据信息

```
In [16]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            1102013 non-null  object
1   TWEET                 1395398 non-null  object
2   STOCK                862231 non-null   object
3   DATE                 862231 non-null   object
4   LAST_PRICE           862231 non-null   float64
5   1_DAY_RETURN         862231 non-null   float64
6   2_DAY_RETURN         862231 non-null   float64
7   3_DAY_RETURN         862231 non-null   float64
8   7_DAY_RETURN         862231 non-null   float64
9   PX_VOLUME            862231 non-null   float64
10  VOLATILITY_10D        862130 non-null   float64
11  VOLATILITY_30D        862184 non-null   float64
12  LSTM_POLARITY         661102 non-null   float64
13  TEXTBLOB_POLARITY     367665 non-null   float64
dtypes: float64(10), object(4)
memory usage: 149.1+ MB
```

```
In [17]: data.head(10)
```

Out[17]:

	Unnamed: 0	TWEET	STOCK	DATE	LAST_PRICE	1_DAY_RETURN	2_DA
0	0	RT @robertoglezcano: @amazon #Patents Show Fl...	NaN	NaN	NaN	NaN	
1	NaN	Amazon	31/01/2017	823.48	0.008379	0.014924	
2	1	@FAME95FM1 Jamaicans make money with @Payoneer...	PayPal	31/01/2017	39.780000	0.002011	
3	2	@CBSi Jamaicans make money with @Payoneer @Pay...	PayPal	31/01/2017	39.780000	0.002011	
4	3	@Hitz92fm Jamaicans make money with @Payoneer ...	PayPal	31/01/2017	39.780000	0.002011	
5	4	RT @loadsofvans: Retweet this post & follo...	NaN	NaN	NaN	NaN	
6	NaN	Amazon	31/01/2017	823.48	0.008379	0.014924	
7	5	RT @Disney: See The Newest Trailer for #Beauty...	NaN	NaN	NaN	NaN	
8	NaN	Disney	31/01/2017	110.65	0.002621	-0.012201	
9	6	RT @tarastrong: Just got @facebook back! I REA...	NaN	NaN	NaN	NaN	

```
In [18]: for i in data.columns:
          if not data[i].isnull().any():
              continue
          data[i] = data[i].fillna(data[i].dropna().mode()[0])
```

填补后的数据集信息如下：

```
In [19]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Unnamed: 0             1395450 non-null object  
1   TWEET                  1395450 non-null object  
2   STOCK                  1395450 non-null object  
3   DATE                   1395450 non-null object  
4   LAST_PRICE             1395450 non-null float64  
5   1_DAY_RETURN           1395450 non-null float64  
6   2_DAY_RETURN           1395450 non-null float64  
7   3_DAY_RETURN           1395450 non-null float64  
8   7_DAY_RETURN           1395450 non-null float64  
9   PX_VOLUME              1395450 non-null float64  
10  VOLATILITY_10D          1395450 non-null float64  
11  VOLATILITY_30D          1395450 non-null float64  
12  LSTM_POLARITY           1395450 non-null float64  
13  TEXTBLOB_POLARITY       1395450 non-null float64  
dtypes: float64(10), object(4)
memory usage: 149.1+ MB
```

In [20]:

data.head(10)

Out[20]:

	Unnamed: 0	TWEET	STOCK	DATE	LAST_PRICE	1_DAY_RETURN	2_DA
0	0	RT @robertoglezcano: @amazon #Patents Show Fl...	Ford	0.0	0.000000	0.000000	
1	Facebook	Amazon	31/01/2017	823.48	0.008379	0.014924	
2	1	@FAME95FM1 Jamaicans make money with @Payoneer...	PayPal	31/01/2017	39.780000	0.002011	
3	2	@CBSi Jamaicans make money with @Payoneer @Pay...	PayPal	31/01/2017	39.780000	0.002011	
4	3	@Hitz92fm Jamaicans make money with @Payoneer ...	PayPal	31/01/2017	39.780000	0.002011	
5	4	RT @loadsofvans: Retweet this post & follo...	Ford	0.0	0.000000	0.000000	
6	Facebook	Amazon	31/01/2017	823.48	0.008379	0.014924	
7	5	RT @Disney: See The Newest Trailer for #Beauty...	Ford	0.0	0.000000	0.000000	
8	Facebook	Disney	31/01/2017	110.65	0.002621	-0.012201	
9	6	RT @tarastrong: Just got @facebook back! I REA...	Ford	0.0	0.000000	0.000000	


比较处理前后数据集差异

下面以1_DAY_RETURN 这一属性为例，对比数据集在剔除含缺失值数据对象后的差异

[illegible][illegible]


The Information of 1 DAY RETURN :

```
count      862231.000000
mean         0.001473
std         0.023068
min        -0.177851
25%        -0.007511
50%         0.000000
75%         0.008192
max         0.243639
Name: 1_DAY_RETURN, dtype: float64
Missing Value Count: 533219
```

[illegible]

The Information of 1 DAY RETURN :

```
count      1.395450e+06
mean       9.103542e-04
std        1.814714e-02
min       -1.778512e-01
25%       -1.564129e-03
50%        0.000000e+00
75%        1.134687e-03
max        2.436389e-01
Name: 1_DAY_RETURN, dtype: float64
Missing Value Count: 0
```

[illegible]

3.4 根据属性的相关关系填补缺失值

我们采用相关系数来判断两个属性间的相似度，并根据属性间的相似度，来寻找可替代的同类型属性。如果相似度低，则用该属性的均值填充。

```
In [22]: data = df.select_dtypes(include=['int64', 'float64']).copy(deep=True)
```

展示填补前的数据信息

```
In [23]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE            862231 non-null float64
1   1_DAY_RETURN          862231 non-null float64
2   2_DAY_RETURN          862231 non-null float64
3   3_DAY_RETURN          862231 non-null float64
4   7_DAY_RETURN          862231 non-null float64
5   PX_VOLUME             862231 non-null float64
6   VOLATILITY_10D        862130 non-null float64
7   VOLATILITY_30D        862184 non-null float64
8   LSTM_POLARITY         661102 non-null float64
9   TEXTBLOB_POLARITY     367665 non-null float64
dtypes: float64(10)
memory usage: 106.5 MB
```

```
In [24]: data.head(10)
```

Out[24]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_VO
0	NaN	NaN	NaN	NaN	NaN	
1	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
4	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
5	NaN	NaN	NaN	NaN	NaN	
6	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	
7	NaN	NaN	NaN	NaN	NaN	
8	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	
9	NaN	NaN	NaN	NaN	NaN	

对于其中所有属性数值都为NaN的数据对象，我们先将这些对象剔除这些对象

```
In [25]: data.dropna(how='all', inplace=True)
```

```
In [26]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 862231 entries, 1 to 1395449
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE             862231 non-null float64
1   1_DAY_RETURN           862231 non-null float64
2   2_DAY_RETURN           862231 non-null float64
3   3_DAY_RETURN           862231 non-null float64
4   7_DAY_RETURN           862231 non-null float64
5   PX_VOLUME              862231 non-null float64
6   VOLATILITY_10D         862130 non-null float64
7   VOLATILITY_30D         862184 non-null float64
8   LSTM_POLARITY          661102 non-null float64
9   TEXTBLOB_POLARITY      367665 non-null float64
dtypes: float64(10)
memory usage: 72.4 MB
```

```
In [27]: data.head(10)
```

Out[27]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_V
1	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
4	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
6	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
8	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	1.2228
10	0.005064	0.014273	0.014273	-0.007290	1.979048e+07	1.5127
11	55.220000	0.012314	0.016298	0.016298	5.831221e-02	1.4307
13	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3028
15	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3028


对于剔除全为NaN的数据行后的数据，我们寻找其中还存在NaN的属性。发现只有VOLATILITY_10D,VOLATILITY_30D,LSTM_POLARITY,TEXTBLOB_POLARITY这四个属性还存在NaN，下面我们对这几个属性进行缺失值填补

```
In [28]: for i in data.columns:
        print("{}:{}".format(i,pd.isnull(data[i]).sum()))
```

LAST_PRICE:0
1_DAY_RETURN:0
2_DAY_RETURN:0
3_DAY_RETURN:0
7_DAY_RETURN:0
PX_VOLUME:0
VOLATILITY_10D:101
VOLATILITY_30D:47
LSTM_POLARITY:201129
TEXTBLOB_POLARITY:494566

下面以热力图的形式展示数据属性间的相关系数

```
In [29]: plt.figure(figsize=(10, 8))
        sns.heatmap(data.corr(), square=True, annot=True)
        plt.show()
```



可以看到PX_VOLUME与VOLATILITY_10D,VOLATILITY_30D相关系数较高,分别为0.3和0.53。VOLATILITY_10D与VOLATILITY_30D相关系数较高为0.53 所以我们尝试用这三个属性互相填补缺失值,对于LSTM_POLARITY和TEXTBLOB_POLARITY这两个属性,由于相关关系并不明显,所以用本身的均值填补


```
In [30]: data["VOLATILITY_10D"] = data["VOLATILITY_10D"].fillna(data["PX_VOLUME"]*0.3 + data["VOLATILITY_30D"] = data["VOLATILITY_30D"].fillna(data["PX_VOLUME"]*0.3 + data
for i in data.columns:
    data[i] = data[i].fillna(data[i].mean())
```

填补后的数据集部分信息如下:

```
In [31]: data.info()

<class 'pandas.core.frame.DataFrame'>
Index: 862231 entries, 1 to 1395449
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE             862231 non-null float64
1   1_DAY_RETURN           862231 non-null float64
2   2_DAY_RETURN           862231 non-null float64
3   3_DAY_RETURN           862231 non-null float64
4   7_DAY_RETURN           862231 non-null float64
5   PX_VOLUME              862231 non-null float64
6   VOLATILITY_10D         862231 non-null float64
7   VOLATILITY_30D         862231 non-null float64
8   LSTM_POLARITY          862231 non-null float64
9   TEXTBLOB_POLARITY      862231 non-null float64
dtypes: float64(10)
memory usage: 72.4 MB
```

```
In [32]: data.head(10)
```

Out[32]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_V
1	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
4	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
6	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
8	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	1.2229
10	0.005064	0.014273	0.014273	-0.007290	1.979048e+07	1.5127
11	55.220000	0.012314	0.016298	0.016298	5.831221e-02	1.4307
13	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3029
15	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3029


比较处理前后数据集差异

下面以VOLATILITY_10D这一属性为例，对比数据集在剔除含缺失值数据对象后的差异

[illegible][illegible]


The Information of VOLATILITY 10D:

```
count      862130.000000
mean        18.293836
std         16.277630
min         -1.000000
25%         8.792000
50%        17.897000
75%        23.916000
max        124.137000
Name: VOLATILITY_10D, dtype: float64
Missing Value Count: 533320
```

[illegible]

The Information of VOLATILITY 10D:

```
count      862231.000000
mean         18.294798
std          16.276925
min          -1.000000
25%           8.801000
50%          17.900000
75%          23.916000
max          124.137000
Name: VOLATILITY_10D, dtype: float64
Missing Value Count: 0
```

[illegible]

3.5 通过数据对象之间的相似性来填补缺失值

我们将数值属性向量化，然后使用K临近算法——KNN来计算数据对象间的距离，以此来判断新数据对象间的相似性。而相似性可以用来判断新数据对象与训练集内数据对象的相似性，从而判断新数据对象的属性。

```
In [34]: numeric_df = df.select_dtypes(include=['int64', 'float64']).copy(deep=True)
```

```
In [35]: numeric_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1395450 entries, 0 to 1395449
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE             862231 non-null float64
1   1_DAY_RETURN           862231 non-null float64
2   2_DAY_RETURN           862231 non-null float64
3   3_DAY_RETURN           862231 non-null float64
4   7_DAY_RETURN           862231 non-null float64
5   PX_VOLUME              862231 non-null float64
6   VOLATILITY_10D         862130 non-null float64
7   VOLATILITY_30D         862184 non-null float64
8   LSTM_POLARITY          661102 non-null float64
9   TEXTBLOB_POLARITY      367665 non-null float64
dtypes: float64(10)
memory usage: 106.5 MB
```

```
In [36]: numeric_df.head(10)
```

Out[36]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_VO
0	NaN	NaN	NaN	NaN	NaN	
1	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
4	39.780000	0.002011	0.012318	0.012318	5.480141e-02	910000
5	NaN	NaN	NaN	NaN	NaN	
6	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	
7	NaN	NaN	NaN	NaN	NaN	
8	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	
9	NaN	NaN	NaN	NaN	NaN	

对于其中所有属性数值都为NaN的数据对象，我们先将这些对象剔除这些对象

```
In [37]: numeric_df.dropna(how='all', inplace=True)
```

In [38]: numeric_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 862231 entries, 1 to 1395449
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE             862231 non-null  float64
1   1_DAY_RETURN           862231 non-null  float64
2   2_DAY_RETURN           862231 non-null  float64
3   3_DAY_RETURN           862231 non-null  float64
4   7_DAY_RETURN           862231 non-null  float64
5   PX_VOLUME              862231 non-null  float64
6   VOLATILITY_10D         862130 non-null  float64
7   VOLATILITY_30D         862184 non-null  float64
8   LSTM_POLARITY          661102 non-null  float64
9   TEXTBLOB_POLARITY      367665 non-null  float64
dtypes: float64(10)
memory usage: 72.4 MB
```

In [39]: numeric_df.head(10)

Out[39]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_V
1	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
4	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.1000
6	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.3447
8	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	1.2228
10	0.005064	0.014273	0.014273	-0.007290	1.979048e+07	1.5127
11	55.220000	0.012314	0.016298	0.016298	5.831221e-02	1.4307
13	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3028
15	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.3028

对于剔除全为NaN的数据行后的数据，我们寻找其中还存在NaN的属性。发现只有VOLATILITY_10D,VOLATILITY_30D,LSTM_POLARITY,TEXTBLOB_POLARITY这四个属性还存在NaN，下面我们对这几个属性进行缺失值填补

```
In [40]: for i in numeric_df.columns:
          print("{}:{}".format(i, pd.isnull(numeric_df[i]).sum()))
```

```
LAST_PRICE:0
1_DAY_RETURN:0
2_DAY_RETURN:0
3_DAY_RETURN:0
7_DAY_RETURN:0
PX_VOLUME:0
VOLATILITY_10D:101
VOLATILITY_30D:47
LSTM_POLARITY:201129
TEXTBLOB_POLARITY:494566
```

调用KNN的包来对数据进行填补，我们选择k=2作为参数 由于knn算法的时间复杂度较大，这里我们仅选择**前50000条数据**进行填补

```
In [41]: data = numeric_df.copy()[0:50000]
```

```
In [42]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 50000 entries, 1 to 84960
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE            50000 non-null   float64
1   1_DAY_RETURN          50000 non-null   float64
2   2_DAY_RETURN          50000 non-null   float64
3   3_DAY_RETURN          50000 non-null   float64
4   7_DAY_RETURN          50000 non-null   float64
5   PX_VOLUME             50000 non-null   float64
6   VOLATILITY_10D        49996 non-null   float64
7   VOLATILITY_30D        49990 non-null   float64
8   LSTM_POLARITY         30854 non-null   float64
9   TEXTBLOB_POLARITY     15041 non-null   float64
dtypes: float64(10)
memory usage: 4.2 MB
```

```
In [43]: from sklearn.impute import KNNImputer
          knn_imputer = KNNImputer(n_neighbors=2)

          df_imputed = knn_imputer.fit_transform(data)

          data = pd.DataFrame(df_imputed, columns=data.columns)
```

展示填补后的数据信息

```
In [44]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LAST_PRICE             50000 non-null   float64
1   1_DAY_RETURN           50000 non-null   float64
2   2_DAY_RETURN           50000 non-null   float64
3   3_DAY_RETURN           50000 non-null   float64
4   7_DAY_RETURN           50000 non-null   float64
5   PX_VOLUME              50000 non-null   float64
6   VOLATILITY_10D         50000 non-null   float64
7   VOLATILITY_30D         50000 non-null   float64
8   LSTM_POLARITY          50000 non-null   float64
9   TEXTBLOB_POLARITY      50000 non-null   float64
dtypes: float64(10)
memory usage: 3.8 MB
```

```
In [45]: data.head(10)
```

Out[45]:

	LAST_PRICE	1_DAY_RETURN	2_DAY_RETURN	3_DAY_RETURN	7_DAY_RETURN	PX_VC
0	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.34470
1	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.10000
2	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.10000
3	39.780000	0.002011	0.012318	0.012318	5.480141e-02	9.10000
4	0.008379	0.014924	0.014924	-0.001263	3.137196e+06	1.34470
5	0.002621	-0.012201	-0.012201	-0.024853	8.485838e+06	1.22290
6	0.005064	0.014273	0.014273	-0.007290	1.979048e+07	1.51210
7	55.220000	0.012314	0.016298	0.016298	5.831221e-02	1.43070
8	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.30290
9	0.010682	0.021363	0.021363	-0.057807	9.469076e+06	3.30290

比较处理前后数据集差异

下面以LSTM_POLARITY这一属性为例，对比数据集在剔除含缺失值数据对象后的差异

总结

至此，我们对Tweet Sentiment's Impact on Stock Returns数据集的预处理和探索性分析全部完成